



Vidur Validation Results for Reviewer 2


(a) Real GPU vs Vidur Prediction



(b) Prediction Error by Batch Size



(c) Linear Model Fit: $\tau = 279.6 + 0.010648 \cdot M$



VALIDATION SUMMARY

=====

Model: Llama-2-7B on NVIDIA A100 80GB

ACCURATE REGION (B = 1-64)

- Samples: 20
- MAPE: 1.61%
- Max Error: 4.12%
- Min Error: 0.03%
- R^2 : 0.9877

Linear Model

- d_0 (fixed overhead): 279.64 ms
- d_1 (per-token cost): 0.010648 ms/token

Assessment: EXCELLENT

All errors < 5% ✓